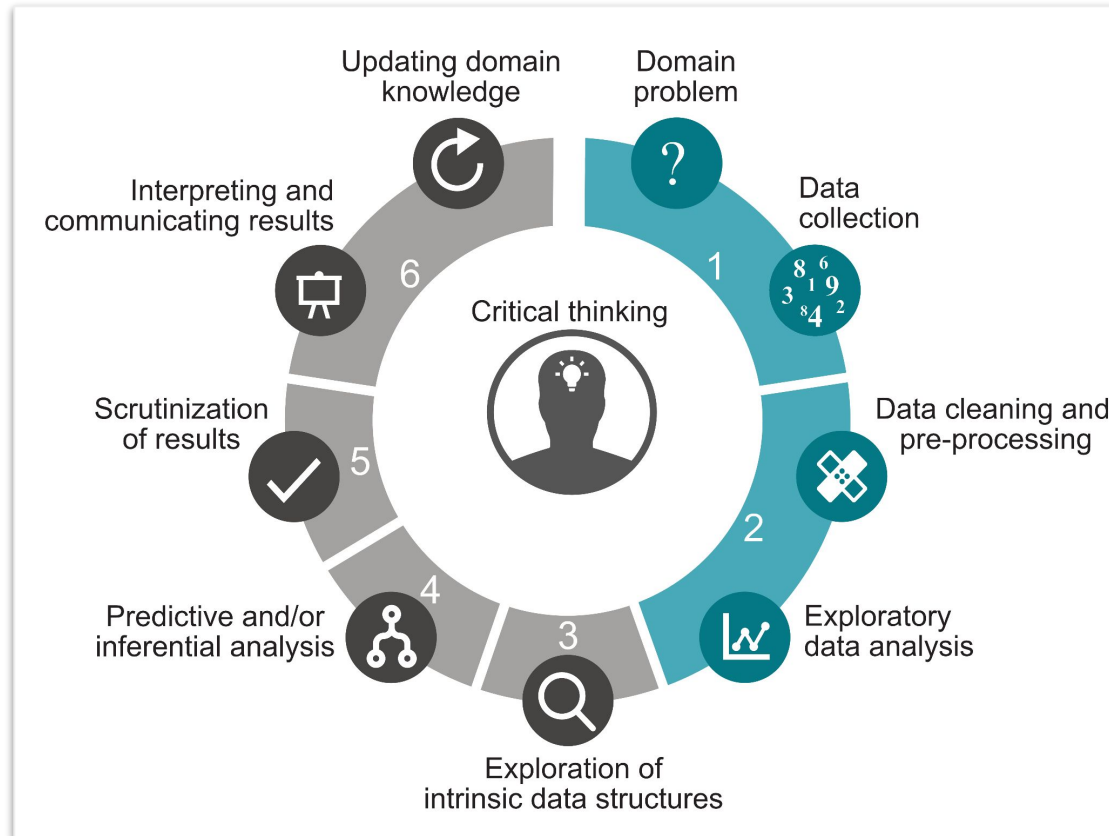


# Introduction to Unsupervised Learning

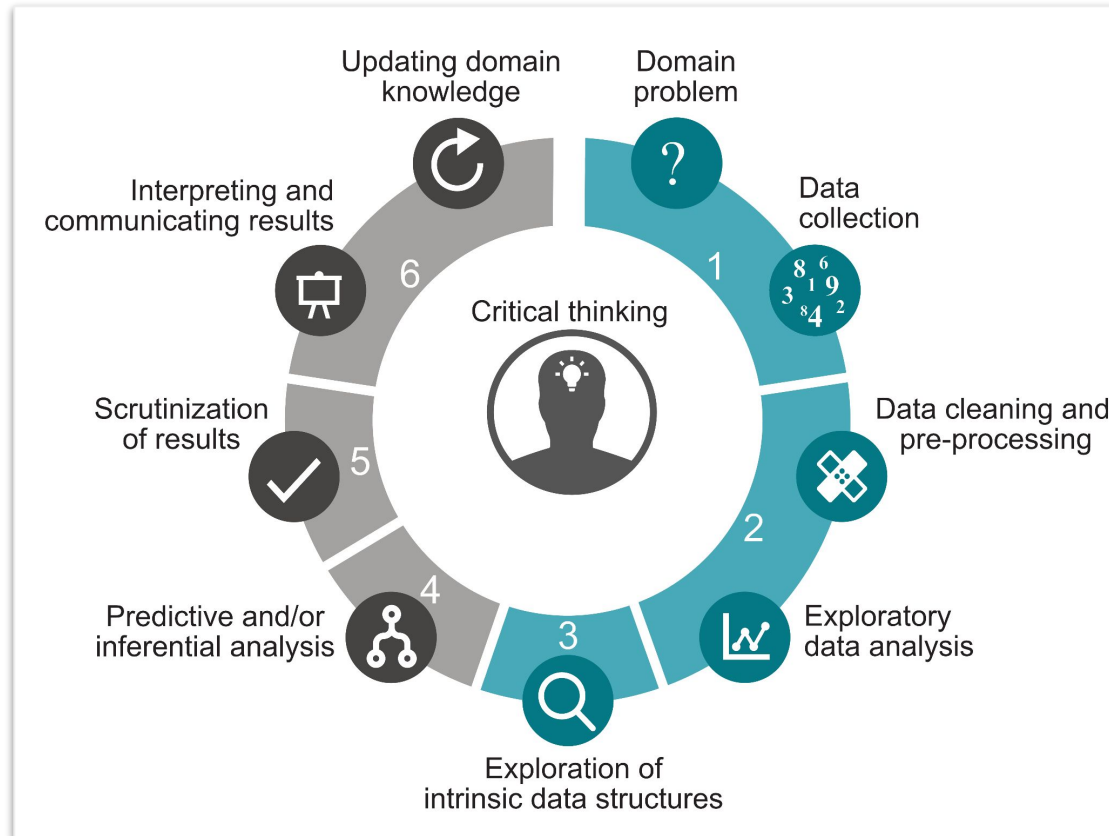
---

February 10, 2025

# The Big Picture: Data Science Life Cycle



# The Big Picture: Data Science Life Cycle



# Today's plan: Introduction to Unsupervised Learning

---


- 1 What is **unsupervised learning**?
- 2 **Applications** of unsupervised learning?
- 3 Overview of popular **dimension reduction** methods
- 4 Overview of popular **clustering** methods

# Supervised vs Unsupervised Learning

In the **supervised learning** setting, we typically have some *covariate/feature* data matrix  $X \in \mathbb{R}^{n \times p}$  and want to predict a *label/response*  $y \in \mathbb{R}^n$

**Example:**


*Spotify data*



# Supervised vs Unsupervised Learning

In the **unsupervised learning** setting, we have some *covariate/feature* data matrix  $X \in \mathbb{R}^{n \times p}$  and but **no label/response**  $y \in \mathbb{R}^n$


**Example:**  
*Spotify data*



# Applications of Unsupervised Learning

In the **unsupervised learning** setting, we have some *covariate/feature* data matrix  $X \in \mathbb{R}^{n \times p}$  and but **no label/response**  $y \in \mathbb{R}^n$

**Example:**  
*Spotify data*



What can we do without labels/responses?

- + **Descriptive statistics**
  - + Ex. mean age of spotify users
- + **Pattern recognition:** discover patterns among observations and/or features
  - + Ex. popular song/genre mashups
- + **Clustering:** identify groups of similar observations and/or features
  - + Ex. groups of people with similar listening histories

\* similar applications in Netflix, Amazon, Youtube, Tiktok, and ad recommendation systems

# Applications of Unsupervised Learning

In the **unsupervised learning** setting, we have some *covariate/feature* data matrix  $X \in \mathbb{R}^{n \times p}$  and but **no label/response**  $y \in \mathbb{R}^n$

**Example:**  
*Political  
Behavior*

**Voter Characteristics**  
(demographics, SES, previous voting behavior, survey data, ...)

**Voters**


$X$

## Clustering/pattern recognition applications:

- + Can we identify groups of similar voters so that we can create targeted messages?
- + Who are the swing voters? And what issues are most likely to sway them?



# Applications of Unsupervised Learning

In the **unsupervised learning** setting, we have some *covariate/feature* data matrix  $X \in \mathbb{R}^{n \times p}$  and but **no label/response**  $y \in \mathbb{R}^n$

**Example:**  
*Anomaly/  
Fraud  
Detection*

Insurance Claims	Features (demographics, billing info, ...)					
			$X$			


## Clustering/pattern recognition applications:

- + Are there any anomalies in the claims/billing data? Maybe these are fraudulent.
- + Can perform clustering to identify similar claims that may be billed incorrectly

# Applications of Unsupervised Learning

In the **unsupervised learning** setting, we have some *covariate/feature* data matrix  $X \in \mathbb{R}^{n \times p}$  and but **no label/response**  $y \in \mathbb{R}^n$

**Example:**  
Cancer  
genomics



## Clustering/pattern recognition applications:

- + Are there **subtypes** of patients who have similar tumors? Moreover, are there particular genes that drive these subtypes?
  - + Hope is that these groups can be treated similarly and in a more personalized way than what's done for the whole group → "personalized medicine"

# How do we "learn" from unsupervised data?

Two common unsupervised learning tools


## 1. Dimension Reduction

- For pattern recognition, visualizing your data, data compression

## 2. Clustering

- For identifying groups or clusters in your data

A common dimension reduction + clustering pipeline:



# Dimension Reduction


---

# Dimension Reduction

In reality, data is often **"high-dimensional"** (i.e., has many covariates/features)

How do we visualize data with  $> 3$  features?


**Dimension Reduction:** aims to find a lower-dimensional representation of the data which preserves as much of the original information as possible



# Principal Components Analysis (PCA)



**Principal Components Analysis (PCA):** finds a lower-dimensional representation of the data which preserves as much of the **variance** in the data as possible

- + More specifically, PCA finds a lower-dimensional hyperplane (or orthogonal directions) such that when the data is projected onto the hyperplane, the variance of the data is maximized
- + PCA is a **linear projection**



# When does PCA "work" and when does PCA "not work"?

*Scenario:* (High-dimensional) Gaussian data




✓ This is the ideal scenario for PCA

# When does PCA "work" and when does PCA "not work"?

## Scenario: Correlated Variables


X and Y are highly correlated;  
Z is independent



### PC Loadings:

$$\text{PC1} = 0.7X + 0.7Y + 0.1Z$$

$$\text{PC2} = -0.1X - 0.1Y + 1.0Z$$





PCs typically group correlated variables together






# When does PCA "work" and when does PCA "not work"?

*Scenario: Highly skewed data or data with outliers*



# When does PCA "work" and when does PCA "not work"?



*Scenario: Highly skewed data or data with outliers*



✓ **if** variance is still a meaningful measure of information

# When does PCA "work" and when does PCA "not work"?

*Scenario: Swiss roll*



✗ not great for nonlinear manifolds


# Nonlinear Dimension Reduction Methods

## tSNE

1. Compute Euclidean distance between every pair of points in  $X$
2. Translate these pairwise distances into probability of being neighbors
  - + Large pairwise distance  $\rightarrow$  low probability of being neighbors
3. Find lower-dimensional representation such that

$$\text{Prob}(i \text{ and } j \text{ are neighbors) in original high-dimensional space} \approx \text{Prob}(i \text{ and } j \text{ are neighbors) in new low-dimensional space}$$

Hyperparameter: perplexity

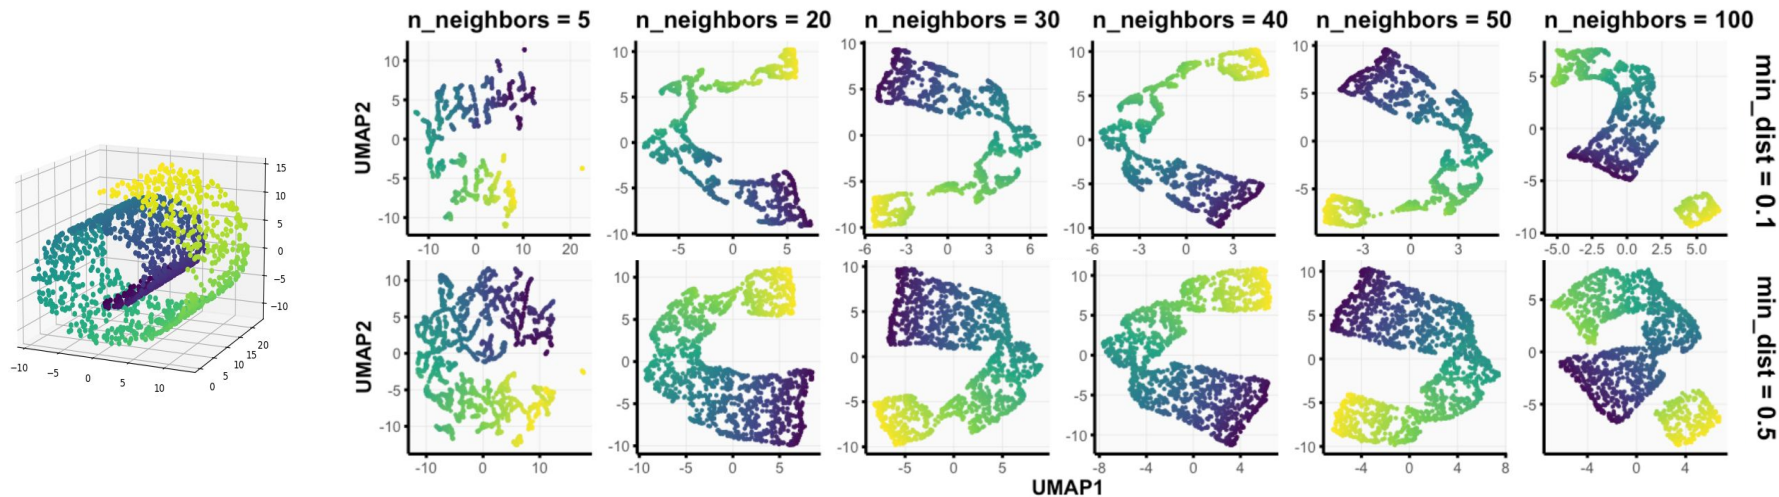


# Nonlinear Dimension Reduction Methods

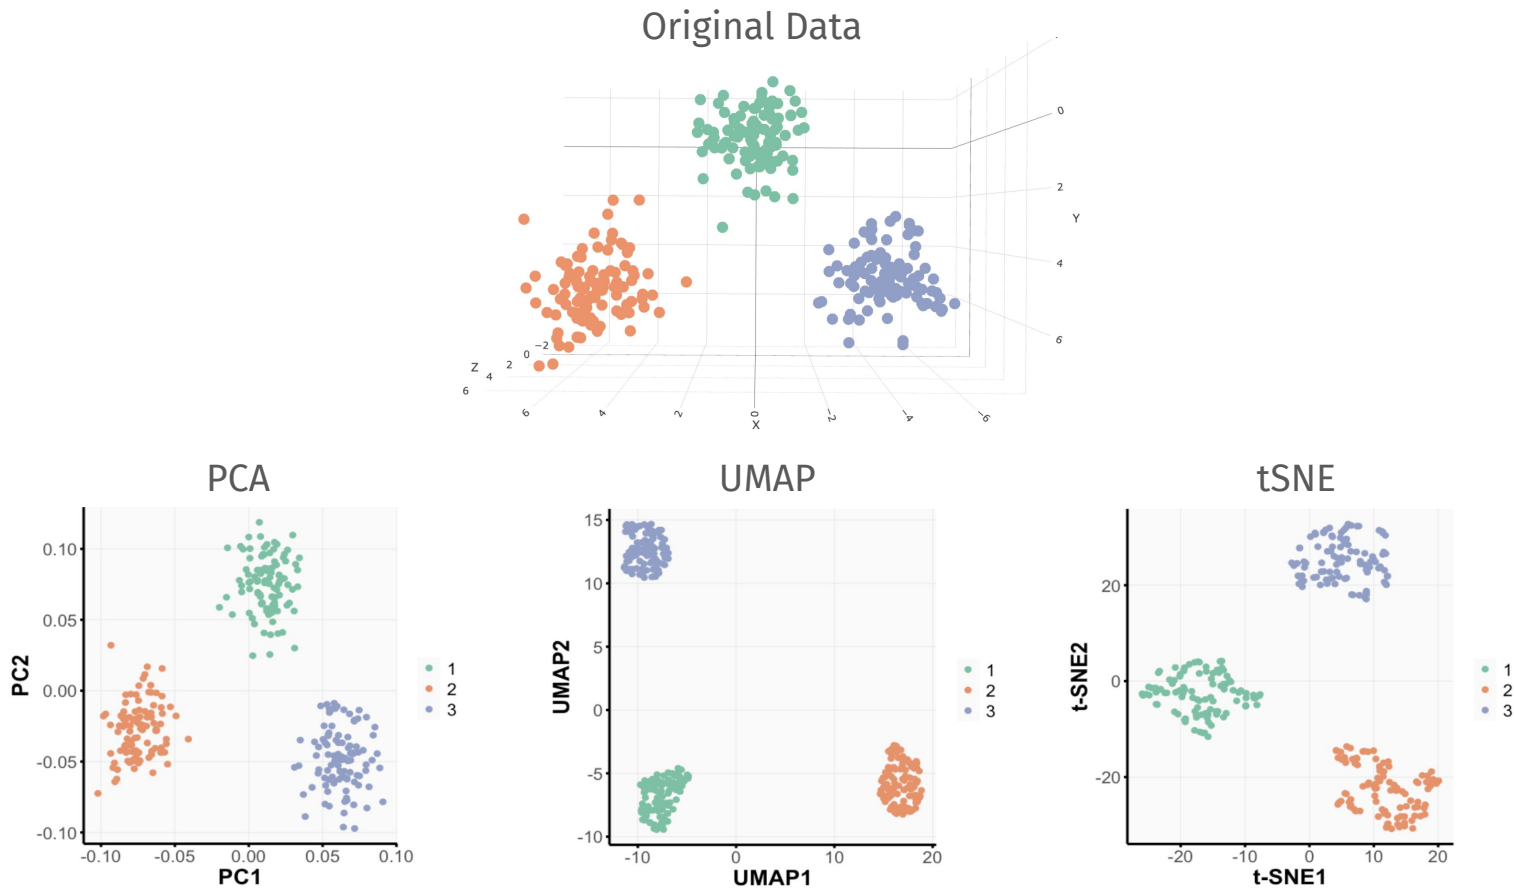
## UMAP

- + UMAP often does better than tSNE at preserving the global structure
- + Like tSNE, the idea is that pairs of points that are close in the original high-dimensional space should also be close in the new low-dimensional space
- + How does UMAP differ from tSNE? Different similarity metrics, loss function, optimization algorithm

Hyperparameter: `n_neighbors` and `min_distance`



# Word of caution: tSNE and UMAP can exaggerate clusters



# Recap: Dimension Reduction Methods

	<b>PCA</b>	<b>tSNE</b>	<b>UMAP</b>
<i>Feature Interpretability</i>	Yes	No	No
<i>Linear/nonlinear</i>	Linear	Nonlinear	Nonlinear
<i>Number of components</i>	Orthogonal, nested; Can compute all $p$ components at once	Non-nested and need to re-run for each chosen rank; Typically only 2-3 components	Non-nested and need to re-run for each chosen rank; Typically only 2-3 components
<i>Computation</i>	Fast	Slower	Slower but faster than tSNE
<i>Unique, global solution</i>	Yes	Converges to local solution	Converges to local solution
<i>Other considerations?</i>	No hyperparameters	Results can change drastically depending on hyperparameters; Not good at preserving global structure; "Curse of dimensionality"	Results can change drastically depending on hyperparameters; Better at preserving global structure than tSNE; "Curse of dimensionality"

Other dimension reduction methods: MDS, NMF, ICA, Isomap, LLE, Autoencoders, ...

# Additional Resources

---

A more detailed review of these unsupervised learning methods + more can be found at <https://tiffanymtang.github.io/dsip-s25/#unsupervised-learning>

- + Also includes R and Python code for implementing these methods

The quarto notebook that generated this walkthrough can be found here:

[https://github.com/tiffanymtang/dsip-s25/blob/main/unsupervised\\_learning/notebooks/unsupervised\\_learning.qmd](https://github.com/tiffanymtang/dsip-s25/blob/main/unsupervised_learning/notebooks/unsupervised_learning.qmd)

## **Additional Resources for Unsupervised Learning Methods**

- + Elements of Statistical Learning Textbook Chapter 14



# Clustering


---

# Clustering

**Clustering:** aims to identify groups/clusters of samples (and/or features) that are "similar"

Two main approaches:


- + K-means clustering
- + Hierarchical clustering



# K-means Clustering

For a pre-specified  $K$ :

- + **Idea:** find  $K$  clusters which result in the "tightest" groups (i.e., has the smallest within-cluster variance)




# K-means Clustering

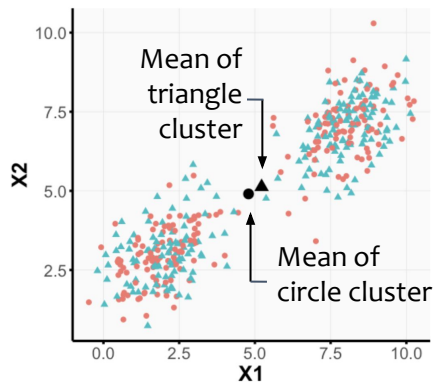
For a pre-specified  $K$ :

- + **Idea:** find  $K$  clusters which result in the "tightest" groups (i.e., has the smallest within-cluster variance)

Original 2D Data




Candidate Cluster A




Blue: high variance  
Red: high variance

Candidate Cluster B



Blue: low variance  
Red: high variance

Candidate Cluster C




Blue: low variance  
Red: low variance

# K-means Clustering


For a pre-specified  $K$ :

- + **Idea:** find  $K$  clusters which result in the "tightest" groups (i.e., has the smallest within-cluster variance)

Original 2D Data




Candidate Cluster A




Blue: high variance  
Red: high variance

Candidate Cluster B



Blue: low variance  
Red: high variance

Candidate Cluster C




Blue: low variance  
Red: low variance

# K-means Clustering

For a pre-specified  $K$ :

- + **Idea:** find  $K$  clusters which result in the "tightest" groups (i.e., has the smallest within-cluster variance)



Points get clustered to the closest centroid

# When does K-means "work" and when does K-means "not work"?


*Scenario:* Spherical, linearly-separable clusters



✓ This is the ideal scenario for K-means

# When does K-means "work" and when does K-means "not work"?

*Scenario: Non-spherical clusters*




✗ Not great for non-spherical clusters



# When does K-means "work" and when does K-means "not work"?

*Scenario: Clusters with different variances*




✗ Not great for clusters with different variances

# Hierarchical Clustering

- + A **greedy, agglomerative** algorithm
- + Gives family of nested clusterings, presented as a tree
- + At the lowest level, each cluster contains a single observation
- + As we move up the tree, some leaves begin to fuse into branches – these are observations that are most **similar** to each other

*Initialization (Step 0):* Each point starts as its own singleton cluster

STEP 0:



DENDROGRAM




# Hierarchical Clustering


- + A **greedy, agglomerative** algorithm
- + Gives family of nested clusterings, presented as a tree
- + At the lowest level, each cluster contains a single observation
- + As we move up the tree, some leaves begin to fuse into branches – these are observations that are most **similar** to each other

*Next Step:* Join the two points/clusters that are "closest" together

STEP 1:



DENDROGRAM




# Hierarchical Clustering


- + A **greedy, agglomerative** algorithm
- + Gives family of nested clusterings, presented as a tree
- + At the lowest level, each cluster contains a single observation
- + As we move up the tree, some leaves begin to fuse into branches – these are observations that are most **similar** to each other

*Next Step:* Join the two points/clusters that are "closest" together

STEP 2:



DENDROGRAM




# Hierarchical Clustering


- + A **greedy, agglomerative** algorithm
- + Gives family of nested clusterings, presented as a tree
- + At the lowest level, each cluster contains a single observation
- + As we move up the tree, some leaves begin to fuse into branches – these are observations that are most **similar** to each other

*Next Step:* Join the two points/clusters that are "closest" together

STEP 3:



DENDROGRAM




# Hierarchical Clustering


- + A **greedy, agglomerative** algorithm
- + Gives family of nested clusterings, presented as a tree
- + At the lowest level, each cluster contains a single observation
- + As we move up the tree, some leaves begin to fuse into branches – these are observations that are most **similar** to each other

*Next Step:* Join the two points/clusters that are "closest" together

STEP 4:



DENDROGRAM




# Hierarchical Clustering


- + A **greedy, agglomerative** algorithm
- + Gives family of nested clusterings, presented as a tree
- + At the lowest level, each cluster contains a single observation
- + As we move up the tree, some leaves begin to fuse into branches – these are observations that are most **similar** to each other

*Next Step:* Join the two points/clusters that are "closest" together

STEP 5:




DENDROGRAM



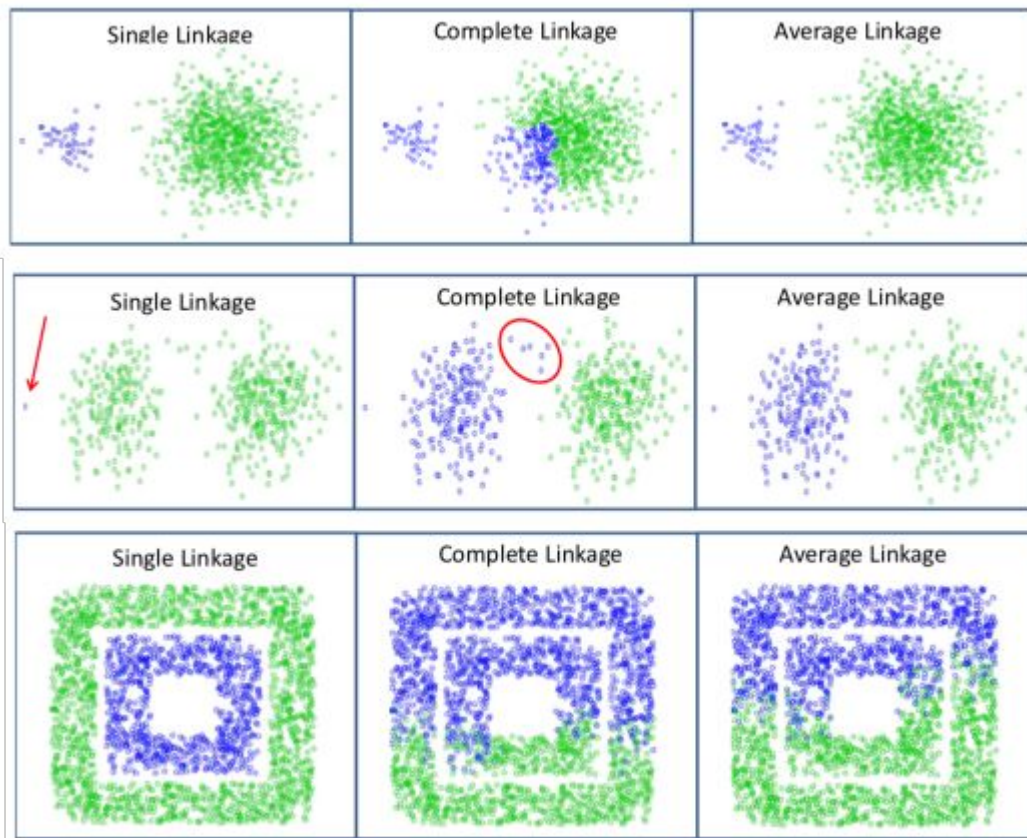
# How to join clusters/observations

1. **Distance metric:** a measure of dissimilarity between two observations
  - a. Examples:  $l_2$ ,  $l_1$ , any of your favorite norms,  $1 - \text{cor}(x, y)$
2. **Linkage metric:** rule for joining two clusters
  - a. Single Linkage (min)
  - b. Complete Linkage (max)
  - c. Average Linkage (average)
  - d. Ward's Linkage (min variance)






# Linkage Examples



# Linkages


## Single Linkage (min)

- + Can handle diverse shapes
- + Very sensitive to outliers or noise
- + Often results in unbalanced clusters
- + Extended, trailing clusters in which observations are fused one at a time – chaining




## Complete Linkage (max)

- + Often gives cluster with similar sizes
- + Less sensitive to outliers
- + Works better with spherical distributions



## Average Linkage

- + Compromise between single & complete linkage
- + Less sensitive to outliers than single linkage, but not as robust as single complete linkage
- + Works better with spherical distributions




**Ward's linkage:** join sets that minimize the Euclidean distance between all pairs of points

**Average and Ward's linkage are the most widely used in practice**

# When does hierarchical clustering "work" or "not work"?


*Scenario: Non-spherical clusters*




# When does hierarchical clustering "work" or "not work"?

## Scenario: Non-spherical clusters


Linkage = single



Linkage = complete




Linkage = average




Linkage = ward.D



Linkage = single




Linkage = complete




X2

Linkage = average




Linkage = ward.D



X1

# When does K-means "work" and when does K-means "not work"?


*Scenario: Clusters with different variances*




# When does K-means "work" and when does K-means "not work"?

*Scenario: Clusters with different variances*


Linkage = single




Linkage = complete




Linkage = average




Linkage = ward.D




Linkage = single




Linkage = complete



Linkage = average




Linkage = ward.D



# Hierarchical Clustering

- + Gives family of nested clusterings, presented as a tree



# Recap: Clustering Methods

	K-means Clustering	Hierarchical Clustering
<i>Advantages</i>	<ul style="list-style-type: none"><li>• Super fast and intuitive</li><li>• Good when clusters are spherical and linearly separable</li></ul>	<ul style="list-style-type: none"><li>• Gives nested family of clusterings</li><li>• Convenient visualizations with dendrograms</li></ul>
<i>Disadvantages</i>	<ul style="list-style-type: none"><li>• Bad when clusters are not spherical or have different variances</li><li>• Must choose K a priori</li><li>• Local solution; depends on initialization</li></ul>	<ul style="list-style-type: none"><li>• Depends <i>heavily</i> on linkage (single, complete, average, Ward's linkage)</li><li>• Greedy search</li></ul>
<i>Shared Disadvantages</i>	<ul style="list-style-type: none"><li>• Irrelevant variables are treated as equals with relevant ones</li><li>• Suffers from "Curse of Dimensionality": computing distances between two points in high dimensions is hard and inaccurate</li></ul>	

→ Do dimension reduction first and then clustering using the dimension-reduced data

Other clustering methods: mixture models, DBSCAN, spectral clustering, K-medoids



# Additional Resources

---

A more detailed review of these unsupervised learning methods + more can be found at <https://tiffanymtang.github.io/dsip-s25/#unsupervised-learning>

- + Also includes R and Python code for implementing these methods

The quarto notebook that generated this walkthrough can be found here:

[https://github.com/tiffanymtang/dsip-s25/blob/main/unsupervised\\_learning/notebooks/unsupervised\\_learning.qmd](https://github.com/tiffanymtang/dsip-s25/blob/main/unsupervised_learning/notebooks/unsupervised_learning.qmd)

## **Additional Resources for Unsupervised Learning Methods**

- + Elements of Statistical Learning Textbook Chapter 14